

clustering

May 12, 2017

1 Part 2 -- Clustering

The libraries that we are going to use:

```
In [14]: import pandas as pd
import scipy as sc
import numpy as np
import nltk
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import adjusted_rand_score
from wordcloud import STOPWORDS
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import ENGLISH_STOP_WORDS
from scipy.spatial.distance import cosine
from nltk.cluster.kmeans import KMeansClusterer
from numpy import array
from __future__ import division
from nltk import cluster
from nltk.cluster import euclidean_distance
from nltk.cluster import cosine_distance
from sklearn.decomposition import TruncatedSVD
from collections import Counter
```

1.1 Setting up

We read our training data:

```
In [15]: mydata = pd.read_csv('train_set.csv', sep='\t')
mydata.head()
```

```
Out[15]:
```

	RowNum	Id	Title \
0	9560	9561	Sam Adams founder: Beer is more than just 'col...
1	10801	10802	Slump in oil prices could mean fall in investm...
2	6726	6727	British Gas owner Centrica warns of higher gas...
3	12365	12366	Ole Gunnar Solskjaer appointed manager of Card...
4	11782	11783	Sunderland target loan signings of Kurt Zouma ...

	Content	Category
0	The craft beer boom, which and been attributed...	Business
1	The International Energy Agency has warned tha...	Business
2	Senior executives at British have been accused...	Business
3	is confident he will have complete control of...	Football
4	Kurt Zouma and Jack Rodwell are on Sunderland...	Football

We declare the stopwords that we are going to use:

```
In [16]: # declaring our stopwords
stopwords = set(STOPWORDS) | set(ENGLISH_STOP_WORDS)
# some additional stopwords based on our own observations
stopwords.add('said')
stopwords.add('say')
stopwords.add('says')
stopwords.add('set')
```

We declare our vectorizer, which is a *TfidfVectorizer* (term-frequency times inverse document-frequency vectorizer), and we pass our data through him:

```
In [17]: vectorizer = TfidfVectorizer(stop_words=stopwords)
X = vectorizer.fit_transform(mydata['Content'])
svd = TruncatedSVD(n_components=100)
X_lsi = svd.fit_transform(X)
vectors = X_lsi
```

1.2 K-Means Clustering

We implement K-Means using Cosine Similarity as a distance function:

```
In [18]: # A IMPLEMENTATION OF CLUSTERING WITH KMEANS USING COSINE SIMILARITY, UTILIZING COMPONENTS
clusterer = cluster.KMeansClusterer(5, cosine_distance, repeats=1)
```

```
In [19]: clusters_array = clusterer.cluster(vectors, True, trace=True) # we take a list of our c
```

```
k-means trial 0
iteration
```

```
/home/marinos/.local/lib/python3.5/site-packages/nltk/cluster/util.py:127: RuntimeWarning: inval
return 1 - (numpy.dot(u, v) / (sqrt(numpy.dot(u, u)) * sqrt(numpy.dot(v, v))))
```

```
iteration
iteration
iteration
iteration
iteration
iteration
```

```
iteration
iteration
iteration
iteration
```

1.3 Printing Clustering Results

We will now work towards printing our Clustering's results in a nice way...

```
In [20]: our_dict = {'0':[], '1':[], '2':[], '3':[], '4':[]}
```

```
    counter = 0
```

```
    # we create a dictionary that for each cluster has the numbers of the texts that belong
    for x in clusters_array :
        our_dict[str(x)] += [counter]
        counter +=1
```

```
In [21]: cnt = Counter()
```

```
    # our categories
```

```
    categories = ['Politics', 'Football', 'Business', 'Technology', 'Film']
```

```
    # we create a counter-dictionary based on the above categories
```

```
    for x in categories :
        cnt[x] = []
```

```
    cnt[''] = []
```

```
    print(cnt)
```

```
Counter({'Politics': [], '': [], 'Technology': [], 'Film': [], 'Football': [], 'Business': []})
```

```
In [22]: data_categories = mydata['Category']
```

```
In [23]: # we will have 5 clusters
```

```
    clusters=['Cluster_0', 'Cluster_1', 'Cluster_2', 'Cluster_3', 'Cluster_4']
```

```
    # we have 5 categories
```

```
    categoryl = ['Politics', 'Business', 'Football', 'Film', 'Technology']
```

```
In [24]: outdict = {'':clusters, 'Technology':[], 'Politics':[], 'Business':[], 'Football':[], 'Film':[]}
```

```
    for cluster_num in range(len(our_dict)): # for each cluster
```

```
        count = {}
```

```
        cluster_length = len(our_dict[str(cluster_num)])
```

```
        for x in our_dict[str(cluster_num)]:
```

```
            category = data_categories[x] # this way we take the category
```

```
            if category in count:
```

```

        count[str(category)] += 1/cluster_length
    else:
        count[str(category)] = 1/cluster_length

    for category in category1: # we create our dictionary
        if str(category) in count:
            outdict[str(category)] += [count[str(category)]]
        else:
            outdict[str(category)] += [0]

print(outdict)

{'': ['Cluster_0', 'Cluster_1', 'Cluster_2', 'Cluster_3', 'Cluster_4'], 'Politics': [0.012581344

```

We create a dataframe with the above data, which we then print to a .csv file:

```

In [25]: # creating the dataframe
out_pd = pd.DataFrame(data=outdict)
out_pd

Out[25]:
   Cluster_0  Cluster_1  Cluster_2  Cluster_3  Cluster_4  Business  Film  Football  Politics  Technology
0  Cluster_0  0.004338  0.960087  0.004772  0.012581  0.018221
1  Cluster_1  0.002374  0.001018  0.993216  0.000000  0.003392
2  Cluster_2  0.082907  0.009086  0.091993  0.016468  0.799546
3  Cluster_3  0.901323  0.001788  0.003218  0.087236  0.006435
4  Cluster_4  0.020774  0.001222  0.004481  0.969857  0.003666

In [26]: # creating the csv file
out_pd = out_pd.ix[:, ['', 'Politics', 'Business', 'Football', 'Film', 'Technology']]
out_pd.to_csv(path_or_buf='clustering_KMeans.csv', sep='\t', index=False)

```